

```
#!/usr/bin/env python3
# -*- coding: utf-8 -*-
"""
@author: fran-pellegrino
"""

# ----- imports -----
import pandas as pd
import numpy as np
import statsmodels.api as sm
import matplotlib.pyplot as plt
from sklearn.model_selection import Ridge, RidgeCV, Lasso, LassoCV, lasso_path, LogisticRegression
from sklearn.preprocessing import train_test_split
from sklearn.preprocessing import StandardScaler
import seaborn as sns
from sklearn.metrics import r2_score, accuracy_score, roc_curve, mean_squared_error, classification_report, roc_auc_score, confusion_matrix
import scipy.stats as stats

# Question 1 below
df = pd.read_csv("techSalaries2017.csv")
cols = df.columns.tolist()
mapped = {
    'company': cols[0],
    'job_title': cols[1],
    'office_location': cols[2],
    'total_annual_compensation': cols[3],
    'base_salary': cols[4],
    'stock_grants': cols[5],
    'bonus': cols[6],
    'years_experience': cols[7],
    'time_with_company': cols[8],
    'gender': cols[9],
    'deg_masters': cols[10],
    'deg_bachelors': cols[11],
    'deg_doctorate': cols[12],
    'deg_highschool': cols[13],
    'deg_somewhere': cols[14],
    'asian': cols[15],
    'white': cols[16],
    'multi_racial': cols[17],
    'black': cols[18],
    'hispanic': cols[19],
    'race_asian': cols[20],
    'education_asian': cols[21],
    'age': cols[22],
    'height_inches': cols[23],
    'zodiac': cols[24],
    'sat': cols[25],
    'gpa': cols[26]
}

numeric_predictors = [
    mapped['years_experience'],
    mapped['time_with_company'],
    mapped['deg_masters'],
    mapped['deg_bachelors'], mapped['deg_somewhere'], mapped['deg_doctorate'],
    mapped['asian'], mapped['white'], mapped['multi_racial'],
    mapped['black'], mapped['hispanic'],
    mapped['age'], mapped['height_inches'],
    mapped['zodiac'], mapped['sat'], mapped['gpa']
]

target = mapped['total_annual_compensation']

# ensure numeric conversion and drop missing
df[target] = pd.to_numeric(df[target], errors='coerce')
for c in numeric_predictors:
    df[c] = pd.to_numeric(df[c], errors='coerce')
data = df[[target] + numeric_predictors].dropna()

# ----- separate X and y -----
X = data.drop(target)
y = data[target]

# ----- multiple regression -----
X_const = sm.add_constant(X)
model_full = sm.OLS(y, X_const).fit()
print("\n==== Multiple Linear Regression Summary ====")
print(model_full.summary())

# full model R^2
r2_full = model_full.rsquared
print(f"\nR^2 for full model: {r2_full:.4f}")

# ----- find best single predictor -----
r2_values = {}

for col in numeric_predictors:
    X_single = sm.add_constant(X[[col]])
    model_single = sm.OLS(y, X_single).fit()
    r2_values[col] = model_single.rsquared

# find best predictor (highest R^2)
best_pred = max(r2_values, key=r2_values.get)
best_r2 = r2_values[best_pred]

print(f"\nBest single predictor: ({best_pred})")
print(f"R^2 (variance explained) by best predictor: {best_r2:.4f}")
print(f"R^2 (variance explained) by full model: {r2_full:.4f}")
print(f"Additional variance explained by other predictors: {r2_full - best_r2:.4f}")

# ----- optional: inspect coefficient for best predictor -----
best_model = sm.OLS(y, sm.add_constant(X[[best_pred]]).fit())
print("\n==== Simple Regression on [best_pred] ====")
print(best_model.summary())

# Visual for predictors ranked by individual R^2
r2_sorted = dict(sorted(r2_values.items(), key=lambda item: item[1], reverse=True))
plt.figure(figsize=(10, 6))
plt.barh(list(r2_sorted.keys()), list(r2_sorted.values()), color='cornflowerblue')
plt.gca().invert_yaxis() # best predictor at top
plt.xlabel('Individual R^2 (Variance Explained)')
plt.title('Predictive Strength of Individual Variables for Total Compensation')
plt.show()

# visual for showing gap between variance explained by experience and variance by full model
# R^2 values
r2_single = 0.1788
r2_full = 0.2857
r2_additional = r2_full - r2_single

# Plot setup
labels = ['Best Predictor\n(Years Experience)', 'Full Model']
values = [r2_single, r2_full]
plt.figure(figsize=(6, 5))
bars = plt.bar(labels, values, color=['skyblue', 'seagreen'])
plt.ylabel('Variance Explained (R^2)')
plt.title('Explained Variance: Best Predictor vs. Full Regression Model')
plt.ylim(0, 0.35)
plt.tight_layout()
plt.show()

# Question 2 below
print("\n==== Ridge Regression Analysis ====")
# ensure numeric conversion and drop missing
df[target] = pd.to_numeric(df[target], errors='coerce')
for c in numeric_predictors:
    df[c] = pd.to_numeric(df[c], errors='coerce')
data = df[[target] + numeric_predictors].dropna()

# ----- X, y split -----
X = data.drop(target)
y = data[target]

# ----- train/test split -----
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# ----- scale predictors -----
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# ----- Ridge regression with CV for alpha -----
alphas = np.logspace(-3, 3, 100)
ridge_cv = RidgeCV(alphas=alphas, store_cv_results=True)
ridge_cv.fit(X_train_scaled, y_train)

best_alpha = ridge_cv.alpha_
print(f"\nBest alpha: {best_alpha:.4f}")

# fit ridge with best alpha
ridge_model = Ridge(alpha=best_alpha)
ridge_model.fit(X_train_scaled, y_train)

# evaluate performance
y_pred_train = ridge_model.predict(X_train_scaled)
y_pred_test = ridge_model.predict(X_test_scaled)

r2_train = r2_score(y_train, y_pred_train)
r2_test = r2_score(y_test, y_pred_test)
rmse_test = np.sqrt(mean_squared_error(y_test, y_pred_test))

print(f"\nTrain R^2: {r2_train:.4f}")
print(f"\nTest R^2: {r2_test:.4f}")
print(f"\nTest RMSE: ${rmse_test:.2f}")

# ----- Coefficients and sparsity info -----
pval_df = Series(ridge_cv.coef_, index=numeric_predictors)
n_total = pval_df.shape[0]
n_zero = (pval_df == 0).sum()
n_nonzero = n_total - n_zero
print(f"\nOut of {n_total} predictors, {n_zero} coefficients were shrunk to EXACTLY zero by Lasso.")
print(f"\nNumber of nonzero coefficients: {n_nonzero}")
print(coef=abs(pval_df).sort_values(ascending=False))

# ----- OPTIONAL: show the selected best single predictor as in prior tasks
# We'll compute single-predictor R^2 for the test set to compare (for chosen best single predictor by coefficient magnitude)
if n_nonzero > 0:
    best_pred = coef.abs().idxmax()
    print(f"\nFeature with largest |beta| in Lasso: ({best_pred})")
    # simple linear regression on scaled single predictor for comparison
    from sklearn.linear_model import LinearRegression
    lr = LinearRegression().fit(X_train_scaled[best_pred], numeric_predictors.index(best_pred).reshape(-1, 1), y_train)
    r2_single_test = r2_score(y_test, lr.predict(X_test_scaled[:, numeric_predictors.index(best_pred)]))

    print(f"\nTest AUC: {auc_B:.4f}, Test Accuracy: {acc_B:.4f}")

# ----- VIUAL 1: Coefficient paths across alphas (lasso_path) -----
# NOTE: We already scaled X_train_s with StandardScaler (mean=0). Do NOT pass fit_intercept to lasso_path.
# Use the same alphas array you fed into Lassocv for consistency (or create a fine grid).
# Compute LASSO path (coefs_path shape: n_features x n_alphas)
alphas_out, coefs_path, _ = lasso_path(X_train_s, y_train, alphas=alphas_path, max_iter=10000)

# Plot coefficient paths
plt.figure(figsize=(12, 8))
for i, feat in enumerate(numeric_predictors):
    plt.plot(np.log10(alphas_out), coefs_path[:, i], label=feat, linewidth=1.25)
plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(feat)
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()

# Question 3 below
print("\n==== Lasso Regression Analysis ====")
# ----- Ensure numeric and drop rows with missing target or predictors -----
df[target] = pd.to_numeric(df[target], errors='coerce')
for c in numeric_predictors:
    df[c] = pd.to_numeric(df[c], errors='coerce')
data = df[[target] + numeric_predictors].dropna().reset_index(drop=True)
print("Rows after dropping NA: ", data.shape[0])

X = data[numeric_predictors].values
y = data[target].values

# ----- train/test split (IMPORTANT: split before any scaling/tuning) -----
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# ----- scale features (Lasso requires scaling for meaningful regularization) -----
scaler = StandardScaler()
X_train_s = scaler.fit_transform(X_train)
X_test_s = scaler.transform(X_test)

# ----- LassoCV to choose alpha (lambda) via cross-validation -----
# Use a log-grid (internally Lassocv will search and return alpha)
alphas = np.logspace(-4, 2, 200)
lasso_cv = LassoCV(alphas=alphas, cv=5, max_iter=10000, n_jobs=1, random_state=42)
lasso_cv.fit(X_train_s, y_train)

best_alpha = lasso_cv.alpha_
print(f"\nOptimal alpha (lambda) found by CV: {best_alpha:.6g}")

# ----- Fit final Lasso model with best alpha -----
lasso_final = Lasso(alpha=best_alpha, max_iter=10000)
lasso_final.fit(X_train_s, y_train)

# ----- Performance on train/test -----
y_pred_train = lasso_final.predict(X_train_s)
y_pred_test = lasso_final.predict(X_test_s)

r2_train = r2_score(y_train, y_pred_train)
r2_test = r2_score(y_test, y_pred_test)
rmse_test = np.sqrt(mean_squared_error(y_test, y_pred_test))

print(f"\nTrain R^2: {r2_train:.4f}")
print(f"\nTest R^2: {r2_test:.4f}")
print(f"\nTest RMSE: ${rmse_test:.2f}")

# ----- Check importance (coefficients) -----
coefs = pd.Series(lasso_final.coef_, index=numeric_predictors)
coefs_sorted = coefs.abs().sort_values(ascending=False)

best_pred = coefs_sorted.index[0]
print(f"\nBest predictor (highest absolute coefficient): {best_pred}")
print(coef=abs(coefs).head(5))

# ----- Visualize coefficient magnitudes -----
plt.figure(figsize=(10, 6))
plt.barh(coefs_sorted.index, coefs_sorted.values, color='teal')
plt.gca().invert_yaxis()
plt.xlabel('Ridge Regression: Predictor Importance (Absolute Coefficients)')
plt.title('Ridge Regression: Predictor Importance (Absolute Coefficients)')
plt.tight_layout()
plt.show()

# ----- Single predictor model for comparison -----
# rebuild ridge with only best predictor
X_best_train = X_train_s[:, best_pred].values
y_best_train = y_train[:, best_pred].values

# Compute LASSO path (coefs_path shape: n_features x n_alphas)
alphas_out, coefs_path, _ = lasso_path(X_train_s, y_train, alphas=alphas_path, max_iter=10000)

# Plot coefficient paths
plt.figure(figsize=(12, 8))
for i, feat in enumerate(numeric_predictors):
    plt.plot(np.log10(alphas_out), coefs_path[:, i], label=feat, linewidth=1.25)
plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(feat)
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()

# Question 4
print("\n==== Logistic Regression Analysis ====")
df_raw = pd.read_csv("techSalaries2017.csv")
cols = df_raw.columns.tolist()
df = df_raw[df_raw['gender'].isin(['Male', 'Female'])].copy()
df['gender_encoded'] = df['gender'].map({'Male': 0, 'Female': 1})

# Ensure numeric conversion for these numeric columns (coerce errors to NaN)
numeric_predictors = [
    mapped['total_annual_compensation'], # compensation (we keep in original units)
    mapped['years_experience'],
    mapped['deg_masters'], mapped['deg_bachelors'], mapped['deg_doctorate'],
    mapped['deg_highschool'], mapped['deg_somewhere'],
    mapped['asian'], mapped['white'], mapped['multi_racial'],
    mapped['black'], mapped['hispanic'],
    mapped['age'], mapped['height_inches'],
    mapped['zodiac'], mapped['sat'], mapped['gpa']
]

for c in numeric_predictors:
    df[c] = pd.to_numeric(df[c], errors='coerce')

# Drop rows with NA in any of the numeric predictors or gender (so both models use the same data)
df_clean = df.dropna(subset=numeric_predictors + ['gender_encoded']).copy()
print("Rows used after filtering and dropping NA: ", df_clean.shape[0])

X = data[numeric_predictors].values
y = data[target].values

# ----- Train/Test split (IMPORTANT: split before any scaling/tuning) -----
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# ----- scale features (Lasso requires scaling for meaningful regularization) -----
scaler = StandardScaler()
X_train_s = scaler.fit_transform(X_train)
X_test_s = scaler.transform(X_test)

# ----- LassoCV to choose alpha (lambda) via cross-validation -----
# Use a log-grid (internally Lassocv will search and return alpha)
alphas = np.logspace(-4, 2, 200)
lasso_cv = LassoCV(alphas=alphas, cv=5, max_iter=10000, n_jobs=1, random_state=42)
lasso_cv.fit(X_train_s, y_train)

best_alpha = lasso_cv.alpha_
print(f"\nOptimal alpha (lambda) found by CV: {best_alpha:.6g}")

# ----- Fit final Lasso model with best alpha -----
lasso_final = Lasso(alpha=best_alpha, max_iter=10000)
lasso_final.fit(X_train_s, y_train)

# ----- Performance on train/test -----
y_pred_train = lasso_final.predict(X_train_s)
y_pred_test = lasso_final.predict(X_test_s)

r2_train = r2_score(y_train, y_pred_train)
r2_test = r2_score(y_test, y_pred_test)
rmse_test = np.sqrt(mean_squared_error(y_test, y_pred_test))

print(f"\nTrain R^2: {r2_train:.4f}")
print(f"\nTest R^2: {r2_test:.4f}")
print(f"\nTest RMSE: ${rmse_test:.2f}")

# ----- Check importance (coefficients) -----
coefs = pd.Series(lasso_final.coef_, index=numeric_predictors)
coefs_sorted = coefs.abs().sort_values(ascending=False)

best_pred = coefs_sorted.index[0]
print(f"\nBest predictor (highest absolute coefficient): {best_pred}")
print(coef=abs(coefs).head(5))

# ----- Visualize partial dependence style curve -----
# Compute bin centers and label each (so legend will include all predictors)
for i, feat in enumerate(numeric_predictors):
    plt.plot(np.log10(alphas_out), coefs_path[:, i], label=feat, linewidth=1.25)
plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(feat)
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()

# Question 5
print("\n==== Extra Credit Part 1 ====")
variables = ['height', 'total_annual_compensation', 'Age']

for var in variables:
    plt.figure(figsize=(12, 4))
    plt.plot(np.log10(alphas_out), coefs_path[:, var], label=var, linewidth=1.25)
    plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(variables[i])
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()

# Extra Credit Part 2
print("\n==== Extra Credit Part 2 ====")
variables = ['height', 'total_annual_compensation', 'Age']

for var in variables:
    plt.figure(figsize=(12, 4))
    plt.plot(np.log10(alphas_out), coefs_path[:, var], label=var, linewidth=1.25)
    plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(variables[i])
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()

# Extra Credit Part 3
print("\n==== Extra Credit Part 3 ====")
variables = ['height', 'total_annual_compensation', 'Age']

for var in variables:
    plt.figure(figsize=(12, 4))
    plt.plot(np.log10(alphas_out), coefs_path[:, var], label=var, linewidth=1.25)
    plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(variables[i])
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()

# Extra Credit Part 4
print("\n==== Extra Credit Part 4 ====")
variables = ['height', 'total_annual_compensation', 'Age']

for var in variables:
    plt.figure(figsize=(12, 4))
    plt.plot(np.log10(alphas_out), coefs_path[:, var], label=var, linewidth=1.25)
    plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(variables[i])
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()

# Extra Credit Part 5
print("\n==== Extra Credit Part 5 ====")
variables = ['height', 'total_annual_compensation', 'Age']

for var in variables:
    plt.figure(figsize=(12, 4))
    plt.plot(np.log10(alphas_out), coefs_path[:, var], label=var, linewidth=1.25)
    plt.title('Partial Dependence Plot (PDP) for each predictor across different alphas')

# Annotate only the top-3 predictors (by absolute final coefficient in the fitted Lasso model)
for i in range(3):
    idx = numeric_predictors.index(variables[i])
    # find closest column in alphas_out to best_alpha
    idx_alpha = np.argmax(np.abs(alphas_out[:, best_alpha]))
    coef_at_best = coefs_path[idx, idx_alpha]
    # place the text little to the right of the vertical line to avoid overlap
    x_text = np.log10(alphas_out[:, best_alpha]) + 0.03
    y_text = coefs_path[i, best_alpha]
    plt.text(x_text, y_text, f"Coef: {coef_at_best:.2f}", ha='center', va='bottom', fontweight='bold',
             bbox=dict(facecolor='white', alpha=0.6, edgecolor='none', pad=1))

# Legend: show full mapping, place it outside the main axes to avoid covering paths
plt.legend(loc='center left', bbox_to_anchor=(1.02, 0.5), fontsize='small', ncol=2)

plt.xlabel('Log10(alpha)')
plt.ylabel('Coefficient Value')
plt.title('Lasso coefficient paths (all predictors) - top 3 labeled on plot')
plt.show()
```